

DSA4211 Report

Aiman Aminuddin

November 2021

Executive Summary

	LASSO	Elastic Net	PCA	PLS	XgBoost
Test MSE	24.25108	24.28296	40.33388	25.82683	25.14059
Test R^2	0.607644	0.6071282	0.347442	0.5821501	0.5932528

Before making any analysis, we realize that one variable: X_{55} with more than 80% of the values missing. We decide to omit X_{55} from our analysis as using only 20% of the values and replacing the NAs with the mean of remaining values is not representative of the distribution of X_{55} . We split the data into a training set to build our models and test set to evaluate the predictive accuracy of said models as well. In order to find the best model, we decide to implement a broad range of approaches to cover all grounds.

The table above shows a summary of the test performance of the different models implemented. It seems that LASSO is the best performing model on test data. This is surprising for several reasons. Theoretically, we would expect Elastic Net to perform better than LASSO as it uses both the l_1 norm and l_2 norm to regularizes the regression model. Having both penalty terms, we have a balance between LASSO and Ridge Regression and perform both variable selection and coefficient reduction (but not exactly 0). On a side note, it is almost as good as LASSO as the difference between test R^2 value is about 10^{-5} . Additionally, XGBoost is a well-regarded as a state-of-the-art algorithm in training models with high accuracy. Many winning entries of Machine Learning Competitions such as Kaggle utilised Xgboost. However here, it is beaten by LASSO.

Additionally, Principal Component Analysis (PCA) is a well-regarded dimension reduction method that takes a large set of predictors and map this set to a smaller set of Principal Components that tries to explain most of the variance in the data set. The idea is that the first M components is sufficient to explain most of the variability in the data set. Here, PCA failed to do so. Using 10-fold Cross-Validation, R suggested that the model with 99 components reduces the MSE the most during training. Clearly, using 99 principal components would lead to over fitting. Thus, it comes to no surprise that the PCA model perform horribly on test data. What worse is that vanilla linear regression model outperforms the PCA model on test data which is surprising. However, this remedied using Partial Least Squares (PLS). Partial Least Squares find directions that help explains both the response and predictors. This is probably why PLS is almost as good as LASSO on test data.

Personally, i find that using LASSO or Elastic Net would be the better strategy in reducing the dimensions of a data set. Not only would they perform moderately well in terms of prediction accuracy on test data, the outputted model is at least interpretable. If the predictors are standardized, we are able to see the relative effect of each predictor on the response. Additionally, the other approaches feels like using a black box approach as we are not able to interpret the relative effect each of the predictor have on a response. Furthermore, having a exact form is nice especially if the data given in a real world setting. Having an exact functional form, would allow people to gain insights the variables that have a greater impact on the response and this could be useful in policy making, industrial process,etc.

DSA4211 Project

Aiman Aminuddin

November 2021

Data Set

Uploading CSV to R

```
1 Data <- read.csv(file.choose(),header = T)
2 nrow(Data) # 1000 rows
3 dim(Data) # 1000 rows and 101 columns
4 names(Data) # 1 Response Variable Y and 100 predictors
```

There are 1000 observations with 1 response variable and 100 predictors: X_1, \dots, X_{100} .

Data Preprocessing

```
1 func1 <- function(x){
2   sum(is.na(x))}
3 apply(Data,2,func1) # check if there are NAs in variables
4
5   Y    X1    X2    X3    X4    X5    X6    X7    X8    X9    X10   X11   X12   X13   X14   X15
6   0     0     0     0     0     0     0     0     0     0     0     0     0     0     0     0
7  X16   X17   X18   X19   X20   X21   X22   X23   X24   X25   X26   X27   X28   X29   X30   X31
8   0     0     0     0     0     0     0     0     0     0     0     0     0     0     0     0
9  X32   X33   X34   X35   X36   X37   X38   X39   X40   X41   X42   X43   X44   X45   X46   X47
10  0     0     0     0     0     0     0     0     0     0     0     0     0     0     0     0
11 X48   X49   X50   X51   X52   X53   X54   X55   X56   X57   X58   X59   X60   X61   X62   X63
12  0     0     0     0     0     0     0     813   0     0     0     0     0     0     0     0
13 X64   X65   X66   X67   X68   X69   X70   X71   X72   X73   X74   X75   X76   X77   X78   X79
14  0     0     0     0     0     0     0     0     0     0     0     0     0     0     0     0
15 X80   X81   X82   X83   X84   X85   X86   X87   X88   X89   X90   X91   X92   X93   X94   X95
16  0     0     0     0     0     0     0     0     0     0     0     0     0     0     0     0
17 X96   X97   X98   X99  X100
18  0     0     0     0     0
19
20
21 Data <- Data[,-56] # Remove X55
```

Using func1 function, we realized that there only one variable with NAs: X_{55} . Additionally, about 81.3% of the values in X_{55} are missing (NAs). We can either remove X_{55} from our analysis or replace NAs with the mean of remaining values. However, only 20% of the data is available for X_{55} . Thus, we decided to remove X_{55} from our analysis. Choosing the latter may result in misrepresentation of the distribution of values of X_{55} as we are only using 20% of the values given.

Currently, we have 99 predictors and 1000 observations. We want to train a regression model from Data to predict values from an independent data set with $m = 10000$. To improve prediction accuracy of our model, we should perform variable selection to reduce the number of predictors used in our model. This is because having many predictors increases the flexibility of our model and result in overfitting. While having a more flexible model improves the fit of model onto the training data. Test RSS will be significantly increase from overfitting. Alternatively, we can implement dimensional reduction methods instead that maps our large

set of predictors to a smaller set of predictors and used them in regression to improve prediction accuracy. Doing so, reduces the flexibility of the model as well and reduce the risk of overfitting.

Splitting Data Set into Training and Test Set

```
1 set.seed(4211) # ensures reproducible results
2 x <- model.matrix(Y~.,Data)[,-1]
3 ind <- sample(2,nrow(x),replace = TRUE,prob = c(0.8,0.2))
4 # Split into training, test set
5 x.train <- x[ind == 1,]
6 x.test <- x[ind == 2,]
7 y.train <- Data$Y[ind == 1]
8 y.test <- Data$Y[ind == 2]
9
10 nrow(x.train)
11 [1] 799
12 nrow(x.test)
13 [1] 201
```

We split Data into a Training Set (x.train,y.train) and Test Set (x.test,y.test). We shall first build a few models from the Training Set and evaluate the performance of these models using the Test Set. The best model will be the one with the lowest MSE and highest R^2 value. Doing so, would give a sensing how well our best model will perform in the independent data set.

LASSO Regression

LASSO performs variable selection by

$$\min \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \|\beta\|_1$$

By having the penalty term $\lambda \|\beta\|_1$, LASSO will shrink the coefficients estimates by a constant amount. However, unlike Ridge Regression (which uses l_2 penalty), the l_1 penalty forces some of the estimates to be exactly 0. Thus, some of the variables are omitted from model (essentially variable selection) and reduces the flexibility of the model. Additionally, note that OLS estimates tend to have larger variance and small bias. In this setting, we should reduce the variance by increasing bias to reduce the risk of overfitting. LASSO does this and doing so, allows for sparse models to be generated instead. Note that when $\lambda = 0$, we recover the Least Squares Regression.

LASSO Coefficient Plots

```
1 library(glmnet)
2 grid <- 10^seq(10,-2,length = 100)
3 # having a large range covers the full range of possible models generated from null model to
4 # least squares fit and choosing the best models from these many models
5 lasso.mod <- glmnet(x.train,y.train,alpha = 1,lambda = grid)
6 plot(lasso.mod)
```

Note that an alternative representation of the problem that LASSO solves is

$$\min \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \text{ subject to } \sum_{j=1}^p |\beta_j| \leq s$$

If s is chosen to be small i.e. arbitrarily close to 0, we would expect many coefficient to shrink to exactly 0 as l_1 norm is close to 0. However, by allowing s to increase, there will be lesser restriction on $\|\beta\|_1$ allowing the coefficients of predictor to increase to their least squares estimates.




Figure 1: LASSO Coefficient of Plot

From Figure 1, the colours represent the coefficients of the predictors. Depending on the choice of s (which translate to the l_1 norm), there will be predictors whose coefficients will shrink to exactly 0 resulting only a subset of predictors to be included in the model.

Plotting Cross-Validation Error (CV) as a Function of λ

```
1 cv.out <- cv.glmnet(x.train,y.train,alpha = 1)
2 plot(cv.out)
```

From Figure 2, we plot CV Error as a function of λ using the grid values of λ . Note that `cv.out` uses 10-fold cross-validation by default. Additionally, the `glmnet()` function by default will standardized the variables inputted so that they are on the same scale for ease of analysis. We want to find λ^* that is the global minimizer of CV Error.

Extracting Optimal λ

```
1 bestlam <- cv.out$lambda.min
2 bestlam
3 [1] 0.1885761
```

We can see that $\lambda^* = 0.1885761$.

Calculating Test MSE and R^2 of Best LASSO Model

```
1 lasso.pred <- predict(lasso.mod,s = bestlam,newx= x.test)
2 # Calculating test MSE
3 mean((lasso.pred-y.test)^2)
4 [1] 24.25108
5 # Computing R^2
6 Rsquared <- function(x,y){
7 1-sum((x-y)^2)/sum((y-mean(y))^2)
8 }
```




Figure 2: LASSO's CV Error as Function of λ

```
9 Rsquared(lasso.pred,y.test)
10 [1] 0.607644
```

On test data, LASSO has a MSE of 24.25108 and R^2 value of 0.607644. We can also compute the test MSE and R^2 of Least Squares Fit model to evaluate LASSO performance on test data.

MSE and R^2 values of Least Squares Fit

```
1 # Test MSE and R^2 of Least Squares Fit
2 linear.pred <- predict(lasso.mod,s = 0,newx = x.test)
3 mean((linear.pred - y.test)^2)
4 [1] 25.43509
5 Rsquared(linear.pred,y.test)
6 [1] 0.5884881
```

It seems that LASSO is slightly better than Least Squares in terms of its test MSE and R^2 value.

Extracting the Coefficients of the LASSO Model

```
1 output <- glmnet(x,Data$Y,alpha = 1,lambda = grid)
2 lasso.coef <- predict(output,type = "coefficients",s=bestlam)[1:100,]
3 lasso.coef[lasso.coef!=0]
4
5 (Intercept)          X1          X7          X8          X9
6 0.8096571959 0.0261793966 0.6733397845 1.9377669710 2.6624478746
7          X10          X20          X22          X23          X26
8 3.8599932239 0.1144885555 -0.0005448235 0.0093351488 0.0117606579
9          X27          X28          X31          X33          X34
10 0.0149140655 -0.0032716856 0.0233850438 0.0480531152 0.0656129865
11          X35          X38          X39          X40          X43
12 0.0734976362 0.0276421016 0.1328009169 0.0069146748 0.1990522916
13          X44          X50          X51          X56          X60
14 0.0132704205 0.1000323688 0.0742940045 0.0015127118 -0.0604258734
15          X62          X68          X69          X70          X71
```

16	0.1063943886	0.2583905989	0.0042553768	0.4305858996	0.5828561665
17	X72	X74	X78	X80	X83
18	-0.0015811229	0.1044664155	0.0331283008	-0.6610099225	0.0282436918
19	X84	X85	X86	X87	X88
20	0.0155329768	0.1009850157	0.2420389856	0.0023711004	0.1370035955
21	X91	X94	X95	X97	X99
22	0.3700325776	0.2182178429	0.1348342967	0.0117315379	0.0010943022

It seems that 44 predictors are included in the model with 55 variables being omitted i.e. coefficients have been shrunk to exactly 0.

Elastic Net Regression

Similar to LASSO, Elastic Net minimizes the residual sum of squares with added penalty. However, Elastic Net uses both the l_1, l_2 norm. Thus, it uses the penalty terms from LASSO and Ridge Regression to regularize the input model. Mathematically, Elastic Net solves:

$$\min \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \left[(1 - \alpha) \frac{\|\beta\|_2^2}{2} + \alpha \|\beta\|_1 \right]$$

where $0 \leq \alpha \leq 1$. As Elastic Net uses both penalty terms from LASSO and Ridge Regression, it is a compromise between LASSO and Ridge. Doing so, could potentially allow Elastic Net to reap the advantages of both methods. For example, LASSO performs variable selection unlike Ridge Regression creating sparse models by shrinking some β_j to exactly 0. While Ridge Regression does not shrink any β_j to exactly 0 (does not perform variable selection), it is a known fact that the variance from Ridge Regression is slightly lower than LASSO and the minimum MSE of Ridge is smaller than that of LASSO. With such advantages from both methods, Thus, Elastic Net take advantages of these strengths of both regularization methods. Note that setting $\alpha = 0$ or $\alpha = 1$, we recover Ridge Regression and LASSO respectively.

Plotting Coefficient Plots for different values of α

```

1 alpha0 <- glmnet(x.train,y.train,alpha = 0,lambda = grid)
2 alpha1 <- glmnet(x.train,y.train,alpha = 0.1,lambda = grid)
3 alpha2 <- glmnet(x.train,y.train,alpha = 0.2,lambda = grid)
4 alpha3 <- glmnet(x.train,y.train,alpha = 0.3,lambda = grid)
5 alpha4 <- glmnet(x.train,y.train,alpha = 0.4,lambda = grid)
6 alpha5 <- glmnet(x.train,y.train,alpha = 0.5,lambda = grid)
7 alpha6 <- glmnet(x.train,y.train,alpha = 0.6,lambda = grid)
8 alpha7 <- glmnet(x.train,y.train,alpha = 0.7,lambda = grid)
9 alpha8 <- glmnet(x.train,y.train,alpha = 0.8,lambda = grid)
10 alpha9 <- glmnet(x.train,y.train,alpha = 0.9,lambda = grid)
11
12 par(mfrow = c(2,5))
13 plot(alpha0,sub = "alpha = 0")
14 plot(alpha1,sub = "alpha = 0.1")
15 plot(alpha2,sub = "alpha = 0.2")
16 plot(alpha3,sub = "alpha = 0.3")
17 plot(alpha4,sub = "alpha = 0.4")
18 plot(alpha5,sub = "alpha = 0.5")
19 plot(alpha6,sub = "alpha = 0.6")
20 plot(alpha7,sub = "alpha = 0.7")
21 plot(alpha8,sub = "alpha = 0.8")
22 plot(alpha9,sub = "alpha = 0.9")

```

From Figure 3, it is easy to see that Elastic Net like LASSO does variable selection for different values of α .

Plotting CV as a Function of λ for each α

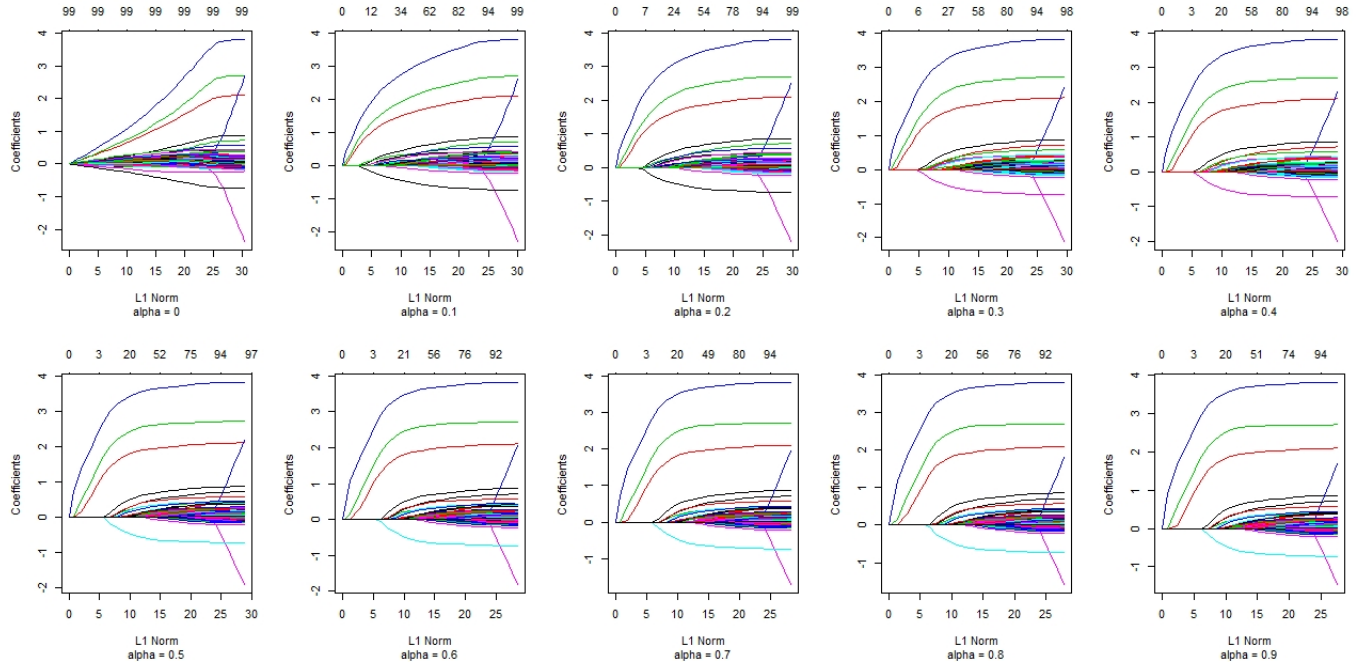


Figure 3: Elastic Net Coefficient Plots for different values of α

```

1 cv.out0 <- cv.glmnet(x.train,y.train,alpha = 0)
2 cv.out1 <- cv.glmnet(x.train,y.train,alpha = 0.1)
3 cv.out2 <- cv.glmnet(x.train,y.train,alpha = 0.2)
4 cv.out3 <- cv.glmnet(x.train,y.train,alpha = 0.3)
5 cv.out4 <- cv.glmnet(x.train,y.train,alpha = 0.4)
6 cv.out5 <- cv.glmnet(x.train,y.train,alpha = 0.5)
7 cv.out6 <- cv.glmnet(x.train,y.train,alpha = 0.6)
8 cv.out7 <- cv.glmnet(x.train,y.train,alpha = 0.7)
9 cv.out8 <- cv.glmnet(x.train,y.train,alpha = 0.8)
10 cv.out9 <- cv.glmnet(x.train,y.train,alpha = 0.9)
11 par(mfrow=c(2,5))
12 plot(cv.out0,sub = "alpha = 0")
13 plot(cv.out1,sub = "alpha = 0.1")
14 plot(cv.out2,sub = "alpha = 0.2")
15 plot(cv.out3,sub = "alpha = 0.3")
16 plot(cv.out4,sub = "alpha = 0.4")
17 plot(cv.out5,sub = "alpha = 0.5")
18 plot(cv.out6,sub = "alpha = 0.6")
19 plot(cv.out7,sub = "alpha = 0.7")
20 plot(cv.out8,sub = "alpha = 0.8")
21 plot(cv.out9,sub = "alpha = 0.9")

```

From Figure 4, it easy to see that for each alpha there exists a λ that is a global minimizer of CV Error. We then extract these λ values.

Extracting optimal λ^* values for each α

```

1 result <- data.frame(alpha = 0,lambda = cv.out0$lambda.min)
2 result <- rbind(result,data.frame(alpha = 0.1,lambda = cv.out1$lambda.min))
3 result <- rbind(result,data.frame(alpha = 0.2,lambda = cv.out2$lambda.min))
4 result <- rbind(result,data.frame(alpha = 0.3,lambda = cv.out3$lambda.min))
5 result <- rbind(result,data.frame(alpha = 0.4,lambda = cv.out4$lambda.min))
6 result <- rbind(result,data.frame(alpha = 0.5,lambda = cv.out5$lambda.min))
7 result <- rbind(result,data.frame(alpha = 0.6,lambda = cv.out6$lambda.min))
8 result <- rbind(result,data.frame(alpha = 0.7,lambda = cv.out7$lambda.min))

```

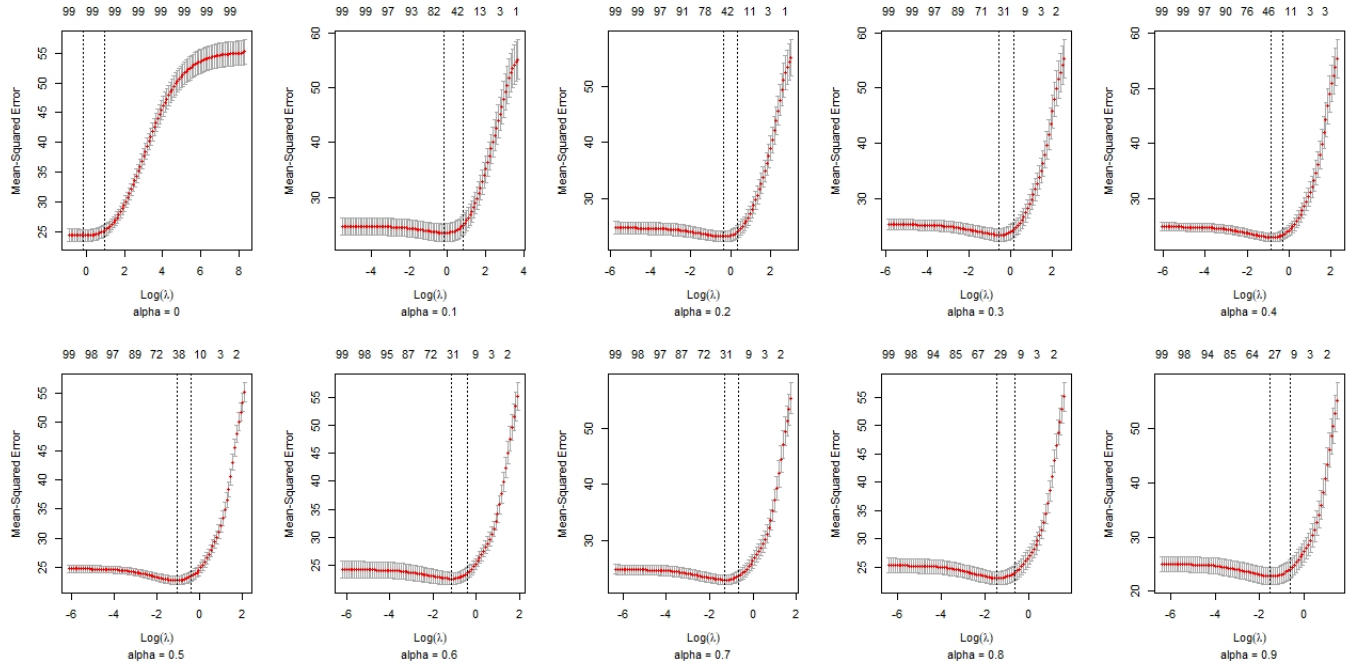


Figure 4: CV Error as a function of λ for each alpha value

```

9 result <- rbind(result,data.frame(alpha = 0.8,lambda = cv.out8$lambda.min))
10 result <- rbind(result,data.frame(alpha = 0.9,lambda = cv.out9$lambda.min))
11 result
12
13   alpha    lambda
14 1     0.0 0.8551698
15 2     0.1 0.8163010
16 3     0.2 0.7132546
17 4     0.3 0.5727451
18 5     0.4 0.4295588
19 6     0.5 0.3436471
20 7     0.6 0.3142935
21 8     0.7 0.2693945
22 9     0.8 0.2357201
23 10    0.9 0.2095290

```

It seems as α increases in value, λ^* decreases in value.

Evaluating Test MSE and R^2 value for different α values

```

1 Elastic.pred0 <- predict(alpha0,s = cv.out0$lambda.min,newx = x.test)
2 Elastic.pred1 <- predict(alpha1,s = cv.out1$lambda.min,newx = x.test)
3 Elastic.pred2 <- predict(alpha2,s = cv.out2$lambda.min,newx = x.test)
4 Elastic.pred3 <- predict(alpha3,s = cv.out3$lambda.min,newx = x.test)
5 Elastic.pred4 <- predict(alpha4,s = cv.out4$lambda.min,newx = x.test)
6 Elastic.pred5 <- predict(alpha5,s = cv.out5$lambda.min,newx = x.test)
7 Elastic.pred6 <- predict(alpha6,s = cv.out6$lambda.min,newx = x.test)
8 Elastic.pred7 <- predict(alpha7,s = cv.out7$lambda.min,newx = x.test)
9 Elastic.pred8 <- predict(alpha8,s = cv.out8$lambda.min,newx = x.test)
10 Elastic.pred9 <- predict(alpha9,s = cv.out9$lambda.min,newx = x.test)
11
12 # Calculating Test MSE
13
14 test0 <- mean((Elastic.pred0-y.test)^2)
15 test1 <- mean((Elastic.pred1-y.test)^2)

```



```

16 test2 <- mean((Elastic.pred2-y.test)^2)
17 test3 <- mean((Elastic.pred3-y.test)^2)
18 test4 <- mean((Elastic.pred4-y.test)^2)
19 test5 <- mean((Elastic.pred5-y.test)^2)
20 test6 <- mean((Elastic.pred6-y.test)^2)
21 test7 <- mean((Elastic.pred7-y.test)^2)
22 test8 <- mean((Elastic.pred8-y.test)^2)
23 test9 <- mean((Elastic.pred9-y.test)^2)
24
25
26 # Calculating test R^2
27
28 R0 <- Rsquared(Elastic.pred0,y.test)
29 R1 <- Rsquared(Elastic.pred1,y.test)
30 R2 <- Rsquared(Elastic.pred2,y.test)
31 R3 <- Rsquared(Elastic.pred3,y.test)
32 R4 <- Rsquared(Elastic.pred4,y.test)
33 R5 <- Rsquared(Elastic.pred5,y.test)
34 R6 <- Rsquared(Elastic.pred6,y.test)
35 R7 <- Rsquared(Elastic.pred7,y.test)
36 R8 <- Rsquared(Elastic.pred8,y.test)
37 R9 <- Rsquared(Elastic.pred9,y.test)
38
39 # Summarizing results
40 test.results <- data.frame(alpha = c(0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9),
41 test.MSE = c(test0,test1,test2,test3,test4,test5,test6,test7,test8,test9),
42 test.Rsquared = c(R0,R1,R2,R3,R4,R5,R6,R7,R8,R9))
43 test.results
44
45      alpha test.MSE test.Rsquared
46 1      0.0 25.50343      0.5873824
47 2      0.1 24.86598      0.5976957
48 3      0.2 24.71573      0.6001266
49 4      0.3 24.61770      0.6017126
50 5      0.4 24.46697      0.6041512
51 6      0.5 24.38123      0.6055383
52 7      0.6 24.36941      0.6057295
53 8      0.7 24.34304      0.6061563
54 9      0.8 24.29137      0.6069921
55 10     0.9 24.28296      0.6071282

```

From test.results, it seems that the best performing model is the model with $\alpha = 0.9$ as it has the smallest test MSE and highest R^2 value.

Extracting the Coefficient of Best Elastic Net Model

```

1 output <- glmnet(x,Data$Y,alpha = 0.9,lambda = grid)
2 Elastic.coef <- predict(output,type = "coefficients",s=cv.out9$lambda.min)[1:100,]
3 Elastic.coef[Elastic.coef!=0]
4
5      (Intercept)          X1          X5          X7          X8
6 0.8089618769 0.0234154893 -0.0000759793 0.6717599407 1.9338277756
7          X9          X10          X13          X20          X22
8 2.6561162337 3.8504479545 -0.0009454334 0.1139382128 -0.0022086818
9          X23          X26          X27          X28          X31
10 0.0034834545 0.0043891085 0.0097817275 -0.0020971076 0.0242015186
11          X33          X34          X35          X38          X39
12 0.0488922351 0.0657789610 0.0753240806 0.0286286455 0.1348755124
13          X40          X43          X44          X50          X51
14 0.0032046096 0.1999437594 0.0111271243 0.0998221465 0.0747285516
15          X56          X57          X60          X62          X65
16 0.0021199463 0.0015740962 -0.0605132441 0.1047029632 0.0012700738
17          X68          X69          X70          X71          X72
18 0.2570226166 0.0028234159 0.4293314364 0.5807641570 -0.0018782354
19          X74          X78          X80          X83          X84
20 0.1053521466 0.0326001853 -0.6602357090 0.0255166394 0.0107869717
21          X85          X86          X87          X88          X91

```




Figure 5: Plot of CV Error as a function of M components

22	0.1014589560	0.2426668174	0.0021958460	0.1359940807	0.3711351162
23	X94	X95	X97	X99	
24	0.2184382665	0.1353811381	0.0046800546	0.0019230803	

Only 48 predictors are included in the model with 51 predictors omitted from it.

Principal Component Analysis

Principal Component Analysis (PCA) is a popular approach to map a large set of predictors to a low dimensional set of features (Principal components). Using PCA, we want to construct the first M principal components to be used as predictors for least squares regression model. The idea is to use a small number of components that is able to explain most of the variability in Data. Doing so, reduces the risk of overfitting.

Plotting CV Error as a function of M

```
1 library(pls)
2 train <- which(ind == 1)
3 pcr.fit <- pcr(Y[,data = Data,subset = train,scale = TRUE,validation = "CV")
4 validationplot(pcr.fit,val.type = "MSEP")
```

Similar to the `glmnet()` function, `pcr` will standardize the predictors inputted for ease of analysis. Furthermore, validation is set to "CV" would mean that 10 fold cross-validation is carried out to determine the optimal M value.

From Figure 5, we see that as M increases, the training MSE decreases. We now extract the optimal value of M.

Extracting the Optimal value of M

```

1 cverr <- RMSEP(pcr.fit)$val[1,,]
2 imin <- which.min(cverr) - 1 # Find the optimal number of PC to use
3 imin
4
5 99 comps
6     99

```

It turns out that to minimize training MSE, we need to use all 99 components. We can also extract the minimum training MSE using the summary function.

Summary of pcr.fit

```

1 summary(pcr.fit)
2
3
4 Data:      X dimension: 799 99
5           Y dimension: 799 1
6 Fit method: svdpc
7 Number of components considered: 99
8
9 VALIDATION: RMSEP
10 Cross-validated using 10 random segments.
11      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
12 CV              7.436  7.441  6.328  6.154  6.107  6.071  6.077
13 adjCV           7.436  7.444  6.046  6.024  6.039  6.041  5.979
14      7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
15 CV          6.043  6.033  6.02  6.037  6.035  6.042  6.014
16 adjCV       5.957  5.974  5.97  6.002  6.011  6.015  5.999
17      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
18 CV          6.004  6.019  6.013  5.986  5.952  5.927  5.936
19 adjCV       5.956  5.982  5.995  5.964  5.905  5.886  5.889
20      21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
21 CV          5.928  5.929  5.911  5.905  5.913  5.913  5.907
22 adjCV       5.892  5.890  5.874  5.879  5.892  5.895  5.862
23      28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
24 CV          5.869  5.869  5.873  5.85  5.842  5.838  5.837
25 adjCV       5.836  5.839  5.849  5.83  5.802  5.804  5.790
26      35 comps 36 comps 37 comps 38 comps 39 comps 40 comps 41 comps
27 CV          5.837  5.840  5.831  5.810  5.821  5.824  5.819
28 adjCV       5.800  5.803  5.807  5.803  5.817  5.805  5.805
29      42 comps 43 comps 44 comps 45 comps 46 comps 47 comps 48 comps
30 CV          5.805  5.809  5.788  5.789  5.782  5.761  5.744
31 adjCV       5.794  5.774  5.761  5.771  5.757  5.713  5.702
32      49 comps 50 comps 51 comps 52 comps 53 comps 54 comps 55 comps
33 CV          5.756  5.759  5.765  5.752  5.734  5.742  5.751
34 adjCV       5.721  5.731  5.729  5.714  5.708  5.722  5.735
35      56 comps 57 comps 58 comps 59 comps 60 comps 61 comps 62 comps
36 CV          5.779  5.802  5.802  5.825  5.798  5.784  5.817
37 adjCV       5.772  5.778  5.784  5.810  5.795  5.791  5.830
38      63 comps 64 comps 65 comps 66 comps 67 comps 68 comps 69 comps
39 CV          5.774  5.762  5.730  5.688  5.682  5.682  5.646
40 adjCV       5.803  5.793  5.772  5.672  5.622  5.626  5.611
41      70 comps 71 comps 72 comps 73 comps 74 comps 75 comps 76 comps
42 CV          5.631  5.608  5.563  5.560  5.531  5.504  5.486
43 adjCV       5.610  5.607  5.590  5.575  5.482  5.447  5.445
44      77 comps 78 comps 79 comps 80 comps 81 comps 82 comps 83 comps
45 CV          5.468  5.486  5.463  5.459  5.456  5.417  5.414
46 adjCV       5.426  5.450  5.435  5.439  5.414  5.391  5.364
47      84 comps 85 comps 86 comps 87 comps 88 comps 89 comps 90 comps
48 CV          5.398  5.363  5.372  5.371  5.329  5.305  5.282
49 adjCV       5.378  5.327  5.347  5.328  5.289  5.292  5.223
50      91 comps 92 comps 93 comps 94 comps 95 comps 96 comps 97 comps
51 CV          5.275  5.254  5.243  5.245  5.237  5.224  5.166
52 adjCV       5.205  5.210  5.207  5.222  5.215  5.211  5.158
53      98 comps 99 comps
54 CV          5.005  4.994
55 adjCV       4.969  4.958

```

```

56
57 TRAINING: % variance explained
58 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
59 X 2.2139 4.081 5.882 7.63 9.289 10.93 12.55 14.16
60 Y 0.9103 36.083 36.163 36.16 36.165 38.48 38.72 38.89
61 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
62 X 15.74 17.30 18.83 20.32 21.80 23.27 24.72
63 Y 38.91 38.97 39.14 39.45 39.54 40.64 40.65
64 16 comps 17 comps 18 comps 19 comps 20 comps 21 comps 22 comps
65 X 26.16 27.57 28.96 30.34 31.69 33.04 34.37
66 Y 40.66 41.14 42.14 42.50 42.71 42.72 42.90
67 23 comps 24 comps 25 comps 26 comps 27 comps 28 comps 29 comps
68 X 35.69 36.98 38.26 39.54 40.80 42.04 43.28
69 Y 43.14 43.14 43.16 43.22 44.12 44.43 44.43
70 30 comps 31 comps 32 comps 33 comps 34 comps 35 comps 36 comps
71 X 44.51 45.71 46.89 48.06 49.21 50.34 51.46
72 Y 44.80 45.08 45.61 45.62 46.13 46.13 46.52
73 37 comps 38 comps 39 comps 40 comps 41 comps 42 comps 43 comps
74 X 52.56 53.66 54.74 55.83 56.89 57.94 58.99
75 Y 46.53 46.55 46.62 47.32 47.38 47.46 48.18
76 44 comps 45 comps 46 comps 47 comps 48 comps 49 comps 50 comps
77 X 60.03 61.06 62.08 63.08 64.08 65.06 66.03
78 Y 48.58 48.67 48.99 49.73 49.91 49.93 50.05
79 51 comps 52 comps 53 comps 54 comps 55 comps 56 comps 57 comps
80 X 66.99 67.92 68.85 69.77 70.67 71.56 72.44
81 Y 50.35 50.53 50.56 50.59 50.66 50.67 51.18
82 58 comps 59 comps 60 comps 61 comps 62 comps 63 comps 64 comps
83 X 73.31 74.18 75.03 75.89 76.73 77.55 78.37
84 Y 51.18 51.23 51.23 51.23 51.31 51.37 51.59
85 65 comps 66 comps 67 comps 68 comps 69 comps 70 comps 71 comps
86 X 79.18 79.98 80.76 81.53 82.30 83.06 83.81
87 Y 51.78 53.64 54.69 54.89 54.91 54.94 54.99
88 72 comps 73 comps 74 comps 75 comps 76 comps 77 comps 78 comps
89 X 84.55 85.27 85.99 86.70 87.40 88.08 88.76
90 Y 55.03 55.51 57.12 57.47 57.59 57.80 57.82
91 79 comps 80 comps 81 comps 82 comps 83 comps 84 comps 85 comps
92 X 89.43 90.09 90.73 91.36 91.98 92.60 93.20
93 Y 58.04 58.08 58.64 58.64 59.22 59.22 60.02
94 86 comps 87 comps 88 comps 89 comps 90 comps 91 comps 92 comps
95 X 93.80 94.39 94.96 95.52 96.07 96.61 97.13
96 Y 60.13 60.75 60.95 60.96 62.15 62.71 62.73
97 93 comps 94 comps 95 comps 96 comps 97 comps 98 comps 99 comps
98 X 97.65 98.16 98.65 99.12 99.56 99.99 100.00
99 Y 62.73 62.75 62.91 62.99 63.54 65.86 66.11
100
101 train.mse <- 4.994 ** 2
102 train.mse
103 [1] 24.94004

```

We have that the training MSE is 24.94004. However, since 99 components are used, there no dimension reduction and we suspect that there is overfitting involved. We now calculate the test MSE and R^2 value of PCR model.

Evaluating the Test MSE and R^2 of PCR Model

```

1 pcr.pred <- predict(pcr.fit,x.test,ncomp=9)
2 mean((pcr.pred-y.test)^2)
3 [1] 40.33388
4 pcr.Rsquared <- Rsquared(pcr.pred,y.test)
5 pcr.Rsquared
6 [1] 0.347442

```

We see as that as compared to Elastic Net and LASSO, PCR perform poorly on test data (even as worse than Least Squares Fit model as well) This could be as a result of overfitting as it uses all the components to fit the training data.

Partial Least Squares

An alternative dimension reduction method is Partial Least Squares (PLS). Similar to PCR, it finds a new set of features: Z_1, \dots, Z_m which are linear combinations of the original predictors. Unlike PCR, PLS utilises the response variable to construct the new features.

Training PLS Model

```
1 pls.fit <- plsr(Y~., data = Data, subset = train, scale = TRUE, validation = "CV")
2 summary(pls.fit)
3
4 Data:      X dimension: 799 99
5           Y dimension: 799 1
6 Fit method: kernelpls
7 Number of components considered: 99
8
9 VALIDATION: RMSEP
10 Cross-validated using 10 random segments.
11      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
12 CV           7.436    5.210    5.050    5.021    5.010    5.002    5.001
13 adjCV        7.436    5.179    5.009    4.984    4.974    4.967    4.966
14      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
15 CV           5.003    4.998    5.000    4.999    5.002    5.001    5.000
16 adjCV        4.968    4.963    4.962    4.962    4.966    4.965    4.965
17      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
18 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
19 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
20      21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
21 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
22 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
23      28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
24 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
25 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
26      35 comps 36 comps 37 comps 38 comps 39 comps 40 comps 41 comps
27 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
28 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
29      42 comps 43 comps 44 comps 45 comps 46 comps 47 comps 48 comps
30 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
31 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
32      49 comps 50 comps 51 comps 52 comps 53 comps 54 comps 55 comps
33 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
34 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
35      56 comps 57 comps 58 comps 59 comps 60 comps 61 comps 62 comps
36 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
37 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
38      63 comps 64 comps 65 comps 66 comps 67 comps 68 comps 69 comps
39 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
40 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
41      70 comps 71 comps 72 comps 73 comps 74 comps 75 comps 76 comps
42 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
43 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
44      77 comps 78 comps 79 comps 80 comps 81 comps 82 comps 83 comps
45 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
46 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
47      84 comps 85 comps 86 comps 87 comps 88 comps 89 comps 90 comps
48 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
49 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
50      91 comps 92 comps 93 comps 94 comps 95 comps 96 comps 97 comps
51 CV           5.000    5.000    5.000    5.000    5.000    5.000    5.000
52 adjCV        4.965    4.965    4.965    4.965    4.965    4.965    4.965
53      98 comps 99 comps
54 CV           5.000    5.000
55 adjCV        4.965    4.965
56
57 TRAINING: % variance explained
58      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
59 X       1.751    2.815    4.219    5.697    6.826    7.921    8.993    9.909
```

60	Y	58.438	64.853	65.659	65.826	65.864	65.873	65.880	65.916
61		9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
62	X	10.52	11.37	12.63	13.75	14.86	15.86	16.87	
63	Y	66.04	66.10	66.11	66.11	66.11	66.11	66.11	
64		16 comps	17 comps	18 comps	19 comps	20 comps	21 comps	22 comps	
65	X	17.94	18.96	20.04	21.22	22.29	23.42	24.55	
66	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
67		23 comps	24 comps	25 comps	26 comps	27 comps	28 comps	29 comps	
68	X	25.56	26.64	27.73	28.78	29.81	30.88	31.98	
69	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
70		30 comps	31 comps	32 comps	33 comps	34 comps	35 comps	36 comps	
71	X	33.05	34.16	35.27	36.32	37.37	38.42	39.48	
72	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
73		37 comps	38 comps	39 comps	40 comps	41 comps	42 comps	43 comps	
74	X	40.58	41.56	42.63	43.71	44.76	45.81	46.80	
75	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
76		44 comps	45 comps	46 comps	47 comps	48 comps	49 comps	50 comps	
77	X	47.78	48.82	49.82	50.86	51.80	52.79	53.77	
78	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
79		51 comps	52 comps	53 comps	54 comps	55 comps	56 comps	57 comps	
80	X	54.83	55.93	56.84	57.78	58.79	59.88	60.87	
81	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
82		58 comps	59 comps	60 comps	61 comps	62 comps	63 comps	64 comps	
83	X	61.82	62.79	63.71	64.65	65.59	66.52	67.46	
84	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
85		65 comps	66 comps	67 comps	68 comps	69 comps	70 comps	71 comps	
86	X	68.39	69.32	70.25	71.19	72.12	73.06	73.99	
87	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
88		72 comps	73 comps	74 comps	75 comps	76 comps	77 comps	78 comps	
89	X	74.92	75.86	76.79	77.72	78.66	79.59	80.53	
90	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
91		79 comps	80 comps	81 comps	82 comps	83 comps	84 comps	85 comps	
92	X	81.46	82.39	83.33	84.26	85.20	86.13	87.07	
93	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
94		86 comps	87 comps	88 comps	89 comps	90 comps	91 comps	92 comps	
95	X	88.00	88.93	89.87	90.80	91.74	92.67	93.60	
96	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	
97		93 comps	94 comps	95 comps	96 comps	97 comps	98 comps	99 comps	
98	X	94.54	95.47	96.41	97.34	98.27	99.21	100.14	
99	Y	66.11	66.11	66.11	66.11	66.11	66.11	66.11	

From inspecting the output of the summary function, it seems that we do not need to use all 99 components in our model to minimize CV Error. Note that by default, `pls()` uses 10 fold cross-validation as well. For easier visualisation let plot CV Error as a function of PLS Components.

Plotting CV Error as a Function of PLS Components

```
1 validationplot(pls.fit, val.type="MSEP")
```

From Figure 6, it seems that we do need to use all 99 components to minimize CV Error. We now extract the optimal number of PLS components from `pls.fit`.

Extracting Optimal number of PLS Components

```
1 cverr <- RMSEP(pls.fit)$val[1,,]
2 imin <- which.min(cverr) - 1
3 imin
4
5 8 comps
6      8
```

Thus, we use 8 PLS Components to minimize CV Error.

Evaluating the Test MSE and R^2




Figure 6: CV Error as a function of PLS Components

```

1 pls.pred <- predict(pls.fit,x.test,ncomp = 8)
2 mean((pls.pred-y.test)^2)
3 [1] 25.82683
4 pls.Rsquared <- Rsquared(pls.pred,y.test)
5 pls.Rsquared
6 [1] 0.5821501

```

PLS does perform better than PCR on test data. However, Elastic Net and LASSO still performs better than PLS.

XGBoost

Extreme Gradient Boosting (XGBoost) is one of the popular machine learning algorithms and has been known to high prediction accuracy. It is considered state-of-the-art machine learning algorithms. XGBoost is essentially a boosting algorithm that can be used to train models for both classification and regression problems. It is in fact an example of Gradient Boosted Decision Tree Algorithm which builds a model from an ensemble of weaker models (gives poor test MSE and R^2 values) such as decision trees.

A generic Gradient Boosted Algorithm involves a training set $\{(x_i, y_i)\}_{i=1}^n$, differentiable Loss Function $L(y, F(x))$ and M number of iterations. Then,

1. Initialise model with a constant value:

$$F_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$$

2. For $m = 1$ to M,

- (a) Compute pseudo-residuals:

$$r_{im} = - \left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right]_{F(x)=F_{m-1}(x)}, i = 1, \dots, n$$

- (b) Fit a base learner (weak learner eg. tree) closed under $h_m(x)$ to pseudo-residual i.e. train it using training set

- (c) Derive γ_m by solving

$$\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, F_{m-1}(x_i) + \gamma h_m(x_i))$$

- (d) Update the model:

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x)$$

3. Output $F_M(x)$ (Friedman, J. H. (February 1999))

Intuitively, Gradient Boosting uses the gradient of the Loss function to minimize the Objective Value. In each iteration, Gradient Boosting uses the gradient of the loss function to find the direction the model parameter vector should take in order to minimize the residual squared error. This is similar to the Gradient Descent algorithm. However, XGBoost uses the Hessian of the Loss function similar to Newton's method to do. This gives better approximated values and thus better gives better directions to minimize residual squared error. Additionally, it utilizes regularization (l_1, l_2 norm) for the same method. (Rblogger)

Note that the parameters used in XGBoost are:

1. nround: maximum number of iteration used by xgboost to train model
2. γ : Regularization parameter. Larger values reduces the risk of overfitting
3. eta: learning rate of XGBoost (Default rate is 0.3, value usually lies between 0.01 and 0.3)
4. maxdepth: controls the depth of the tree (smaller trees reduce the risk of overfitting)
5. subsample: controls the number of samples given to each tree (learner)
6. colsamplertree: is the proportion of predictors to be used
7. minimum child weight: Parameters is used to minimized possible predictor interaction to reduce risk of overfitting
8. max delta step: refer to the maximum weight given to each decision tree (hackerearth)

Finding optimal parameters for XGBoost Model using 10 fold Cross-Validation

```
1 library(xgboost)
2 d.train <- xgb.DMatrix(x.train, label = y.train)
3 d.test <- xgb.DMatrix(x.test, label = y.test)
4
5 # Iterate through 200 times to find best parameters
6 # to minimize Training MSE
7
8
9 optimal.parameters <- list() # store the optimal parameters after for loop
10 optimal.seed <- 4211
11 optimal.rmse <- Inf # store minimum root MSE
12 optimal.idx <- 0 # Locate the smallest minimum root MSE
13
14 for (iter in 1:200){
15   parameters <- list(objective = "reg:linear",
16     eval_metric = "rmse",
```



```

17     max_depth = sample(2:10,1),
18     eta = runif(1,0.01,0.3),
19     subsample = runif(1,0.6,0.9),
20     colsample_bytree = runif(1,0.5,0.8),
21     min_child_weight = sample(1:40,1),
22     max_delta_step = sample(1:10,1))
23
24 # At each iteration make use of Randomisation to minimize
25 # training MSE
26
27 no.rounds <- 1000
28 cv.folds <- 10 # Performing 10 fold Cross-Validation
29 current.seed <- sample.int(10000,1) # find the best seed that minimizes CV
30 set.seed(current.seed)
31
32 model <- xgb.cv(data = d.train,params = parameters,nfold = cv.folds,
33               nrounds = no.rounds,verbose = F,
34               early_stopping_rounds = 8,maximise = FALSE)
35
36 current.idx <- model$best_iteration
37 current.rmse <- model$evaluation_log[current.idx]$test_rmse_mean
38 # update the optimal rmse
39 if (current.rmse < optimal.rmse){
40     optimal.parameters <- parameters
41     optimal.idx <- current.idx
42     optimal.seed <- current.seed
43     optimal.rmse <- current.rmse}
44
45 }
46
47 # Source: https://yangliuresearch.blogspot.com/2018/07/extreme-gradient-boosting-xgboost.html

```

Optimal Parameters obtained from Random Search

```

1 set.seed(optimal.seed)
2 xg_mod <- xgboost(data = d.train,params = optimal.parameters,nround = optimal.idx,verbose =
  F)
3
4 # Best Tuning Parameters
5 result <- data.frame(optimal.parameters,optimal.idx,optimal.rmse,optimal.seed)
6 result
7
8   objective eval_metric max_depth      eta subsample colsample_bytree
9 1 reg:linear      rmse          2 0.02979434 0.7835935      0.7236754
10 min_child_weight max_delta_step optimal.idx optimal.rmse optimal.seed
11 1              19              3         370      4.837736         575
12

```

Feature Importance

```

1 importance_matrix <- xgb.importance(colnames(x.train), model = xg_mod)
2 library(Ckmeans.1d.dp) # for xgb.ggplot.importance
3 xgb.ggplot.importance(importance_matrix, top_n = 10,
4                       measure = "Gain")
5 # Source: https://yangliuresearch.blogspot.com/2018/07/extreme-gradient-boosting-xgboost.html
6

```

From Figure 7, it seems that X10 is by far the most important feature followed by X9 and 8 (by a significant margin).

Plotting the first few Trees of the XGBoost Model

```

1 library(DiagrammeR)

```




Figure 7: XGBoost Feature Importance

```
2 xgb.plot.tree(model = xg_mod, trees = 1)
3 xgb.plot.tree(model = xg_mod, trees = 2)
4 xgb.plot.tree(model = xg_mod, trees = 3)
```

As mentioned previously, Xgboost is an ensemble of weaker model. We will plot the first few decision trees of the Xgboost model. From Figure 8,9,10, we have the first 3 trees of the Xgboost model.

Test MSE and R^2 value

```
1 # Calculate test MSE and R^2
2 yhat_xg <- predict(xg_mod, d.test)
3 test.mse <- mean((yhat_xg - y.test)^2)
4 test.mse
5 [1] 25.14059
6 Rsquared(yhat_xg, y.test)
7 [1] 0.5932528
```

It seems that XGBoost is almost as good as LASSO but LASSO is still the best performing model on test data.

Predicting Unseen Data

```
1 # LASSO is the best performing model on test data
2 # we shall use LASSO to predict values of y from an independent set
3 # Upload Data Set to R
4 test <- read.csv(file.choose(), header = T)
5 test <- test[, -55] # remove X55
6 edited <- data.frame(Y = rep(1, nrow(test)), test)
7 final.test <- model.matrix(Y ~ ., edited)[, -1]
8 Y <- predict(lasso.mod, s = bestlam, newx = final.test)
9 values <- data.frame(Y = Y)
10 library(readr)
11 setwd("C:\\Users\\Aiman\\Documents\\DSA4211")
12 write_csv(values, "A0187108Bfinal.csv")
```

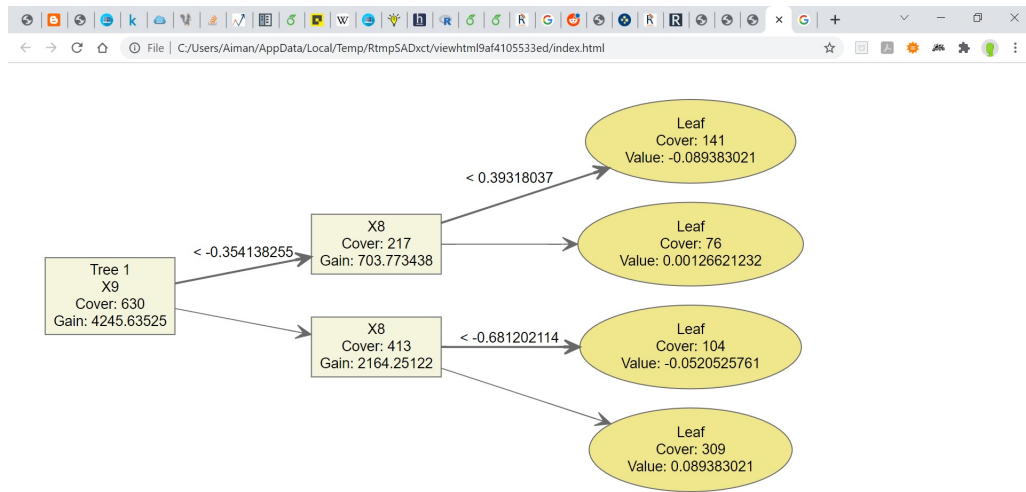


Figure 8: First Tree of the Xgboost model

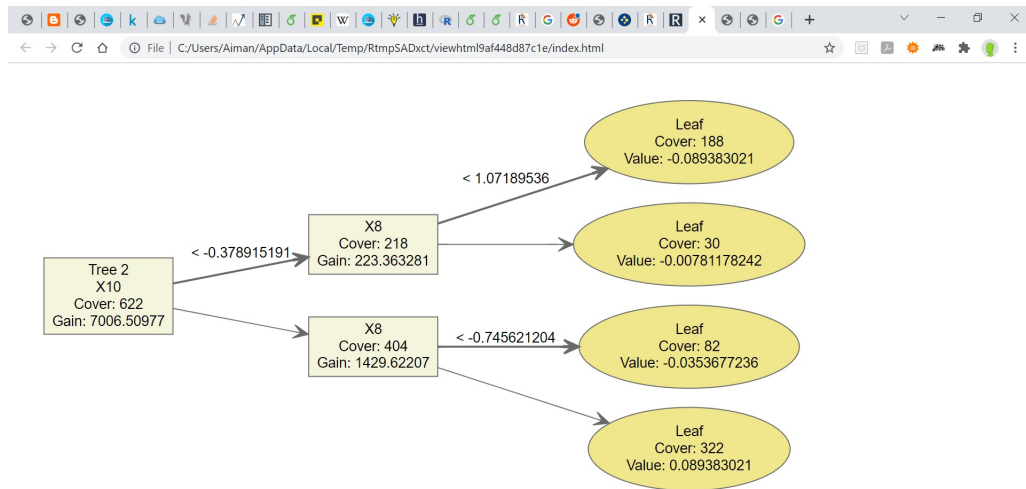


Figure 9: Second Tree of the Xgboost model

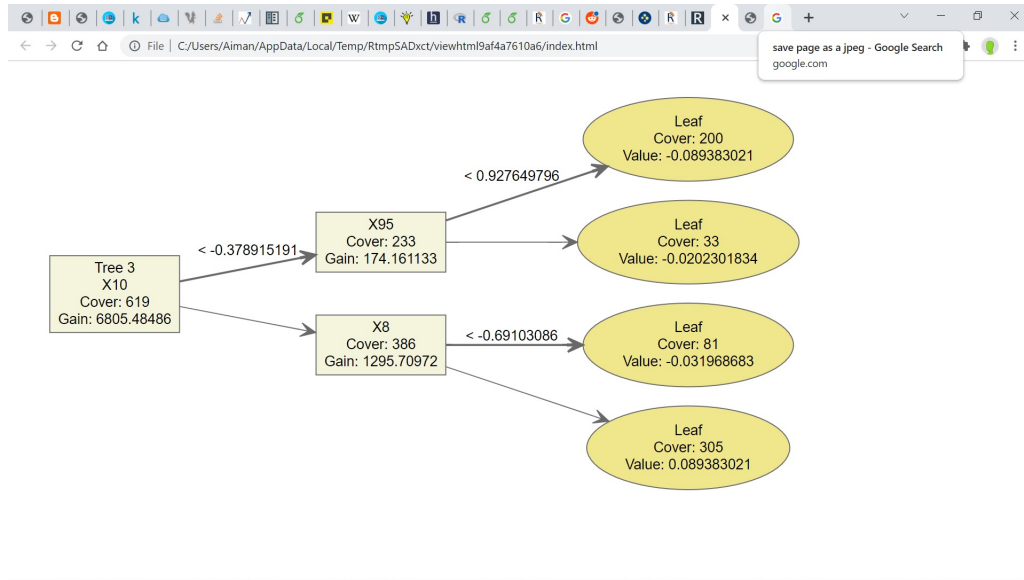


Figure 10: Third Tree of Xgboost Model

References

1. Friedman, J. H. (February 1999). "Greedy Function Approximation: A Gradient Boosting Machine"
2. Beginners tutorial on XGBoost and parameter tuning in R tutorials amp; notes: Machine learning. HackerEarth. (n.d.). Retrieved November 8, 2021, from <https://www.hackerearth.com/practice/machine-learning/machine-learning-algorithms/beginners-tutorial-on-xgboost-parameter-tuning-r/tutorial/>.
3. Glander, D. S. (2018, November 29). Machine learning basics – gradient boosting amp; xgboost: R-bloggers. R. Retrieved November 8, 2021, from <https://www.r-bloggers.com/2018/11/machine-learning-basics-gradient-boosting-xgboost/>.
4. Liu, Y. (2018, July 10). Extreme gradient boosting (XGBoost): Better than random forest or gradient boosting. eXtreme Gradient Boosting (XGBoost): Better than random forest or gradient boosting. Retrieved November 8, 2021, from <https://yangliuresearch.blogspot.com/2018/07/extreme-gradient-boosting-xgboost.html>.